

ECEN 5053-002

Developing the Industrial Internet of Things

Week 1

Expectations and Course Overview

Dave Sluiter - Spring 2018

Expectations and Course Overview

Who Am I?

- 30+ year chip design engineer
- Unisys (Sperry Univac) - Mainframe floating point unit
- Artist Graphics - 3D graphics chips
- LSI Logic -
 - North American Field Coreware Manager - MIPS CPUs, CoreWare Methodology
 - Lead hardware architect for satellite set-top box decoder chips
- Seagate/Micron - Hardware architect for enterprise solid state drives

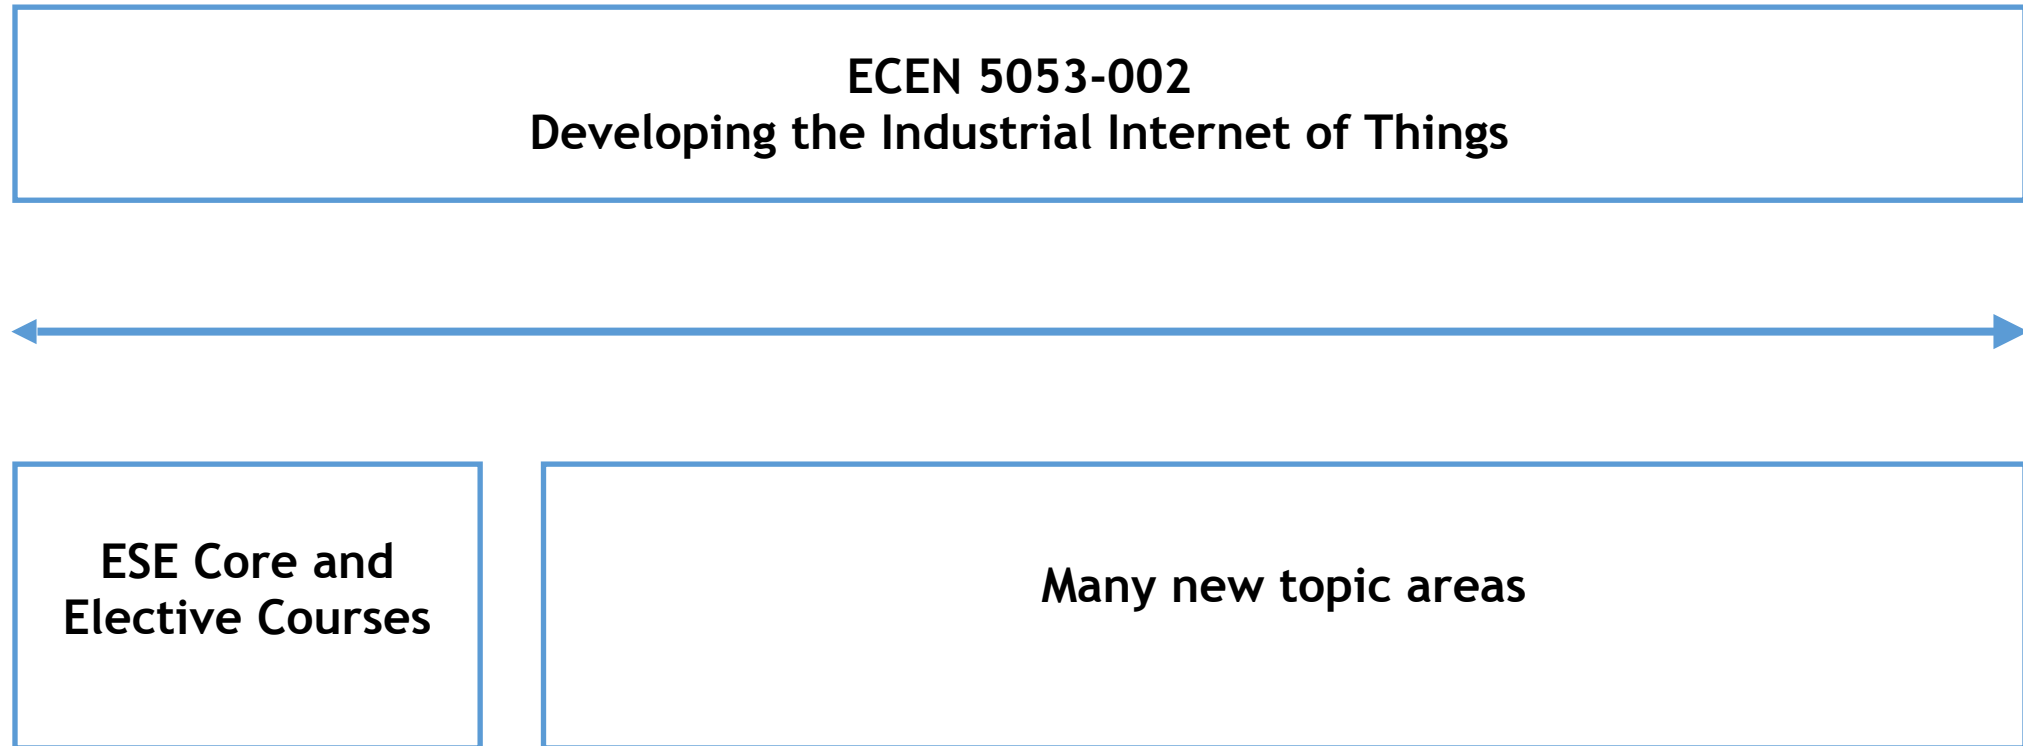
Logistics

- Communication:
 - D2L
 - Email


Expectations

- You are Graduate students!
- Attendance/Participation is important
- Honor Code Violations:
 - **Do not** represent someone else's work as your own
 - **I do want** collaboration and exchange of ideas
 - **I do not want** copying - Do your own work!
 - As graduate students, you will not be given a 2nd chance

Where does this course fit?



Course Overview

- Format
 - Traditional lecture
- Survey Course
 - Intention is to introduce you to a wide set of subject matter
 - Broaden your scope/viewpoint/vocabulary
- Assessments
 - Attendance/Participation
 - Small number of projects - easy, provides hands-on experience
 - Midterm exam
 - Final exam
- No homework 

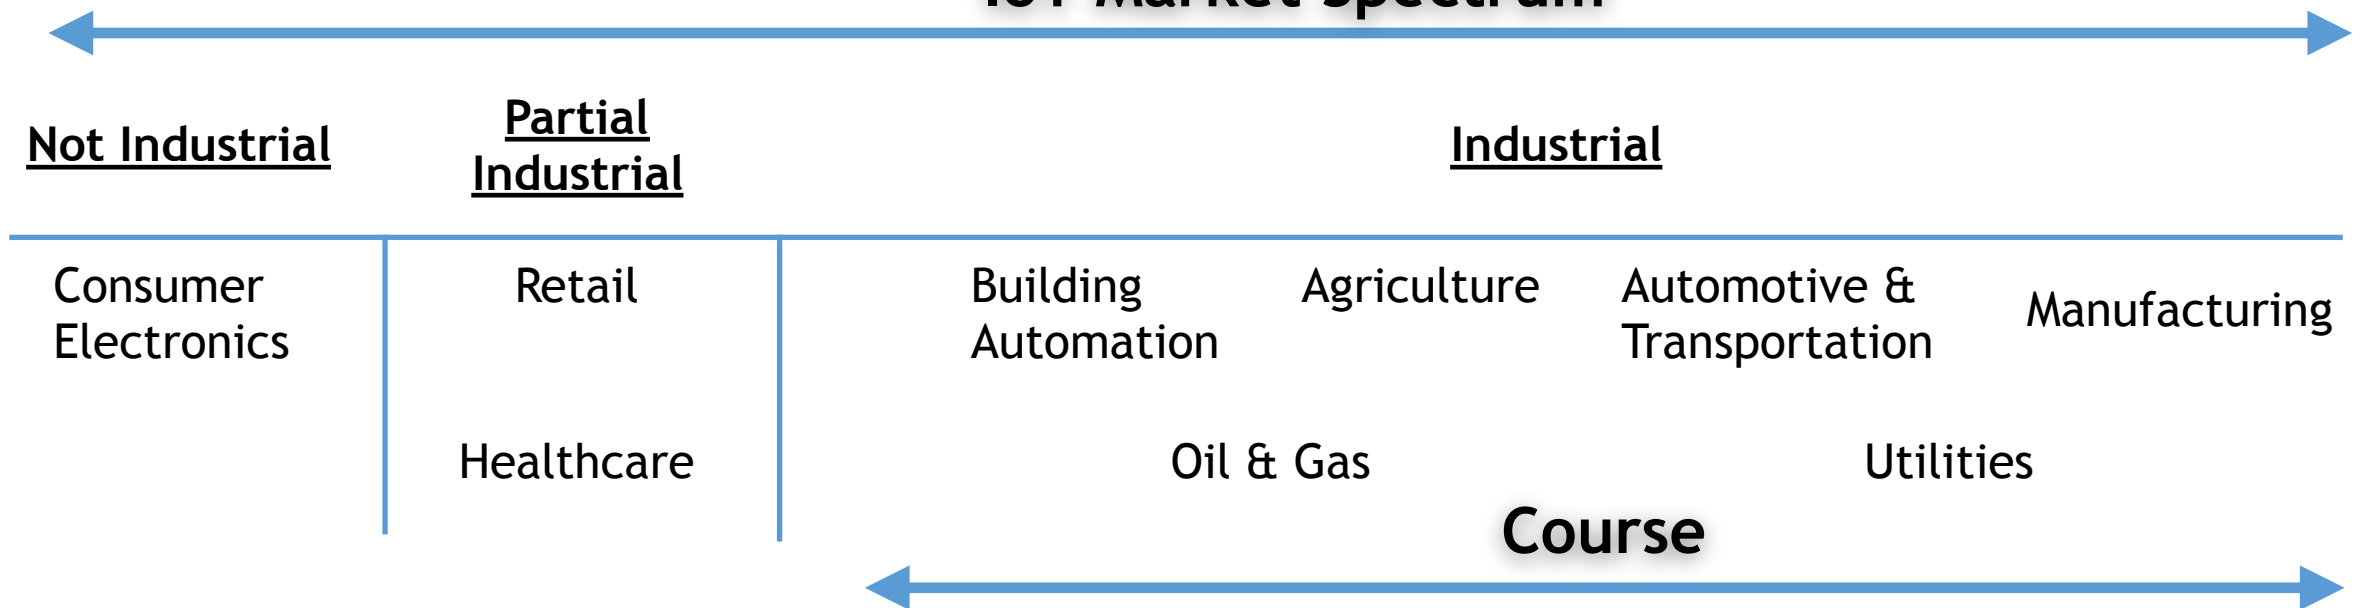
Course Overview (con't)

- Primary focus
 - Emerging trends: Market and Technical
 - Key business concepts for engineers
 - Key skills to develop
 - Understanding the “big picture”, how these systems are built and the value propositions they offer
 - Your role in DIIoT - Show you how your ESE core and elective courses intersect with IIoT

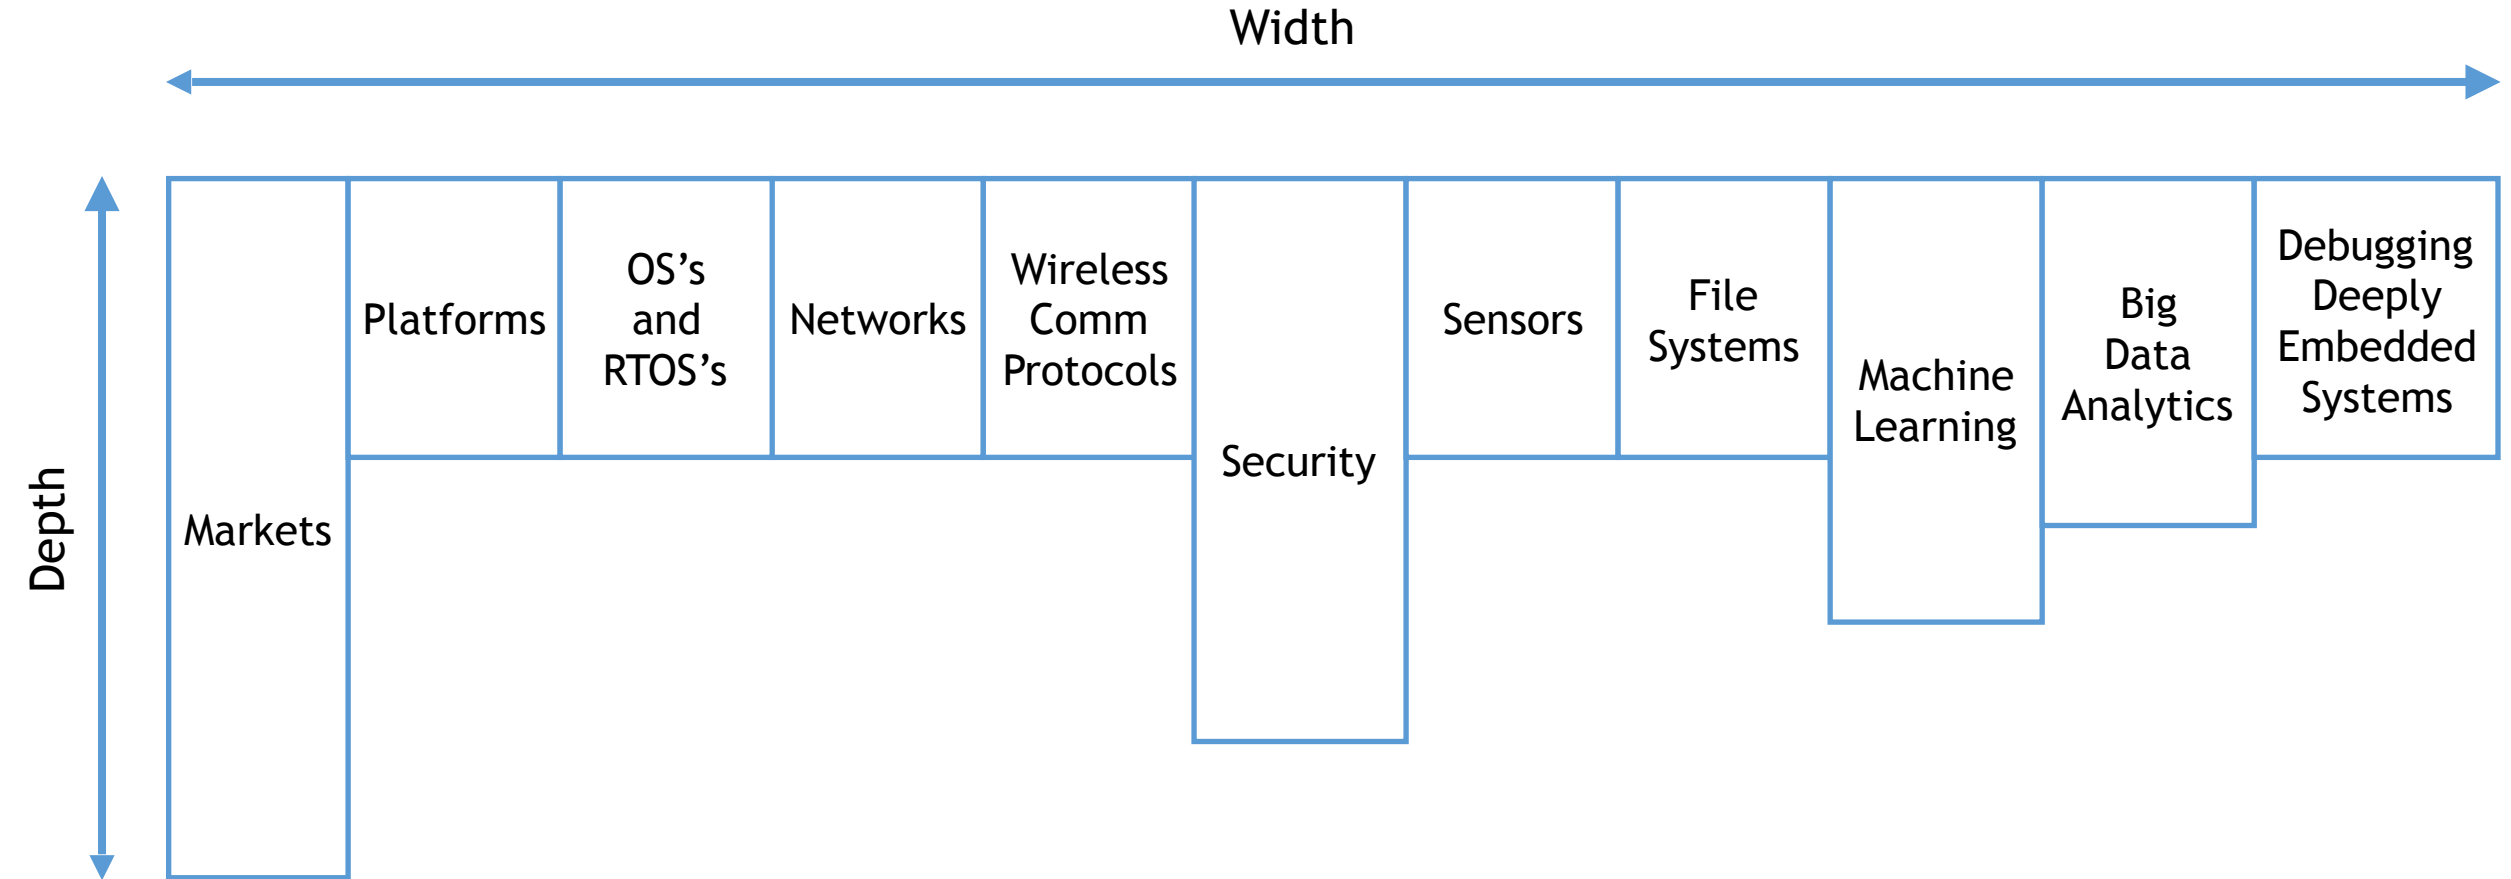
Course Overview (con't)

- What is Industrial IoT?

IoT Market Spectrum





Topic Areas




What will be covered

- An inside look at markets; size and \$ estimates
- Platforms, IBM Bluemix + Watson
- Networks, NFV and SDN
- Security
- Project planning, staffing and execution
- Sensors, file systems
- Machine learning
- Big data analytics
- SystemC
- Debugging deeply embedded systems
- Guest speakers




Security

What does it mean to be secure?




What does it mean to be secure?



Encryption Techniques

- Caesar Cipher, named after Julius Caesar
- Substitution cipher, based on a shift



Encryption Techniques

- One time pad (OTP)
- So-called “perfect” encryption
- Impractical for real-world


A=0, B=1, C=2, ... Z=25											
Plain text	M	E	E	T	T	O	N	I	G	H	T
	12	4	4	19	19	14	13	8	6	7	19
Key	D	Z	H	S	U	I	M	W	E	K	C
	3	25	7	18	20	8	12	22	4	10	2
Sum	15	29	11	37	39	22	25	30	10	17	21
Sum mod 26	15	3	11	11	13	22	25	4	10	17	21
Cipher Text	P	D	L	L	N	W	Z	E	K	R	V
Cipher text	P	D	L	L	N	W	Z	C	K	R	V
	15	3	11	11	13	22	25	4	10	17	21
Key	D	Z	H	S	U	I	M	W	E	K	C
	3	25	7	18	20	8	12	22	4	10	2
Diff	12	-22	4	-7	-7	14	13	-18	6	7	19
Sum mod 26	12	4	4	19	19	14	13	8	6	7	19
Plain text	M	E	E	T	T	O	N	I	G	H	T



AES (Advanced Encryption Standard)


- Established by US NIST (National Institute of Standards)
- Block cipher: 16-bytes in, 16-byte out
- 3 key lengths: 128-, 192-, 256-bits
- High-level description
 - 1) With N =number of rounds, round keys are extracted from the cipher key (where $N = 10, 12, 14$, for 128-, 192- or 256-bits)
 - 2) Round 0
 - 3) Rounds 1 to $N-2$
 - 4) Final round $N-1$
- Believed to be secure, but we don't how to prove it.

Asymmetric Encryption



Diffie-Hellman

A method to securely establish a known secret (a “key”) between 2 parties over an insecure channel.



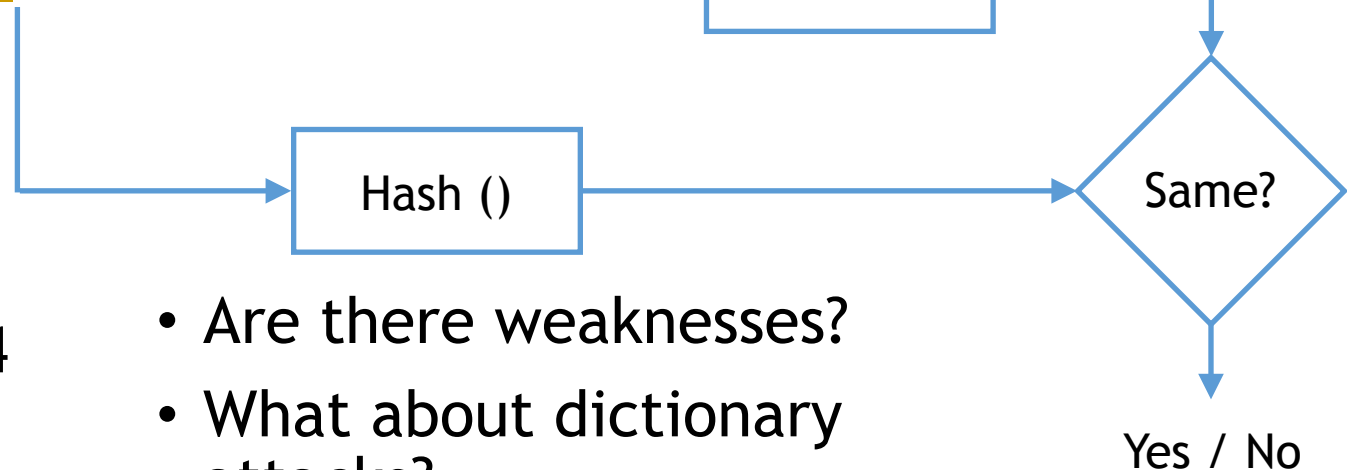
Source: https://en.wikipedia.org/wiki/Diffie-Hellman_key_exchange

Uses for Hash Functions

- Saving passwords
 - Login: **UserName**
 - Password: **Xh8_i27vZ**

Password Table

0xF1324578



Yes / No

- Are there weaknesses?
- What about dictionary attacks?

- See also the SHA-2 family of hash functions as per FIPS 180-4
- Well studied, haven't spotted a problem yet

Take Aways

- Security Mind-set
- Data Integrity
- Authentication
- Encryption
- Will look at few “security blunders”

Machine Learning


Linear Regression

- We have some example data
 - Training data
- Each row is an example from real sales data
- Contains a number of features, x_n
 - x_1 is square footage
 - x_2 is number of bedrooms
- And output y , the price
 - Is what we are trying to predict, also known as the **target value**

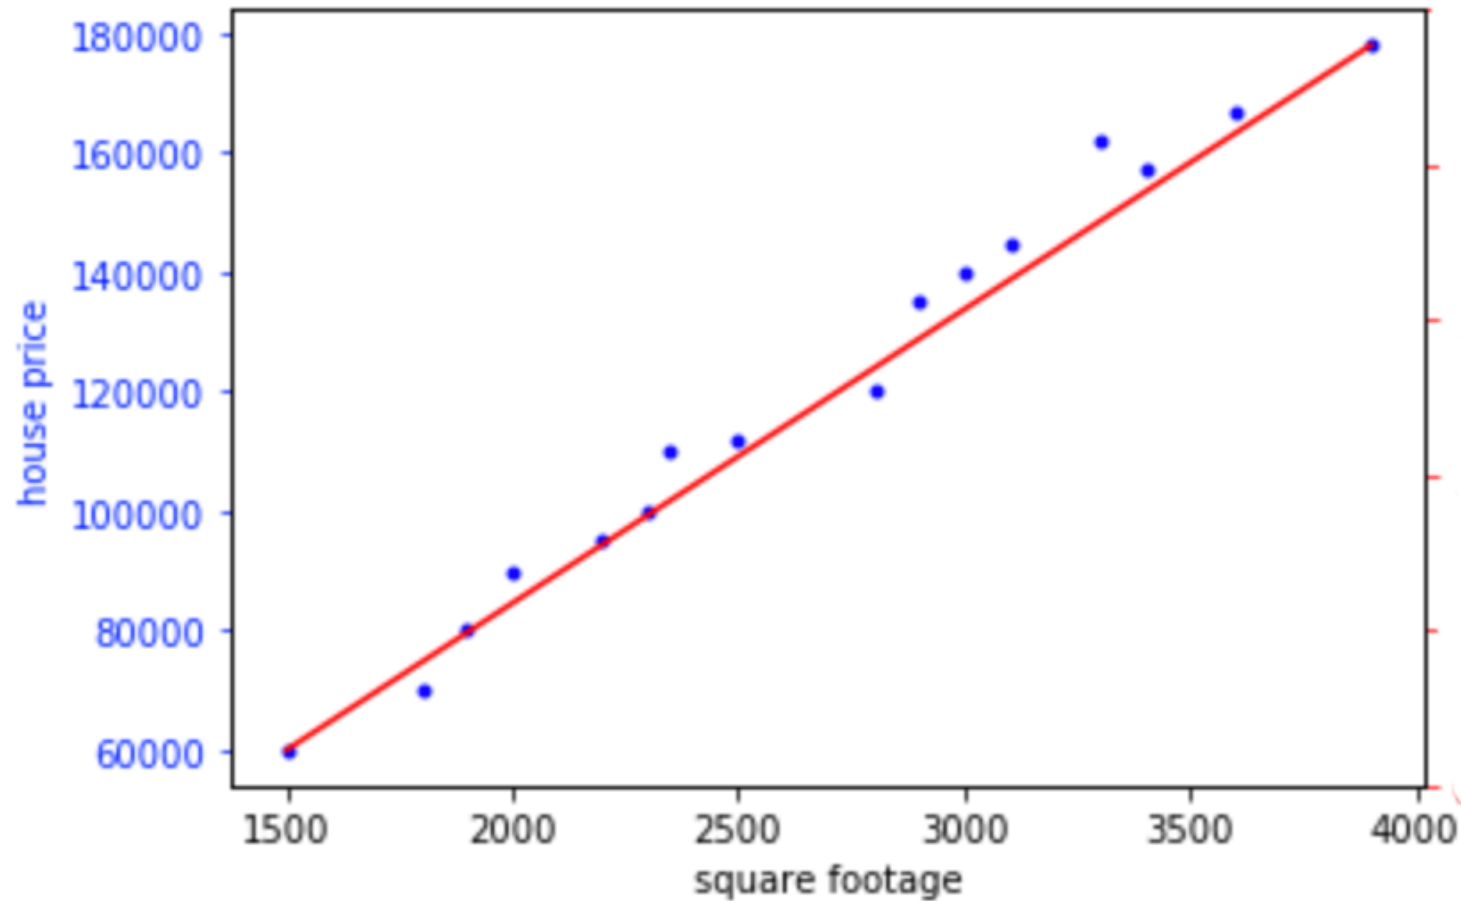
```
train_data = np.array(  
    [  
        # sqft, #bedrooms, price  
        [1500, 2, 60000],  
        [1800, 2, 70000],  
        [1900, 2, 80000],  
        [2000, 3, 90000],  
        [2200, 3, 95000],  
        [2300, 2, 100000],  
        [2350, 3, 110000],  
        [2500, 3, 112000],  
        [2800, 4, 120000],  
        [2900, 3, 135000],  
        [3000, 4, 140000],  
        [3100, 4, 145000],  
        [3300, 5, 162000],  
        [3400, 4, 157000],  
        [3600, 5, 167000],  
        [3900, 5, 178000]  
    ]  
    ) # end training data
```

Training Data

of bedrooms is not plotted



We can manually draw a best-fit curve



Linear Regression

- We want to create a hypothesis function $h()$ that will make predictions
- To perform supervised learning, we have to decide how we will represent the hypothesis function $h()$
- As an initial choice, let's say we decide to approximate $h()$ as a linear combination of features x and weights θ , so we have:
 - $h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2$
 - where we set $x_0 = 1$, and θ_0 becomes the intercept term
 - like in the equation for a line $y = ax + b$, b is the intercept term

Linear Regression

- m = number of training examples, 16
- n = number of features, 2
- We can rewrite $h_{\theta}(x)$ as :

$$\bullet \quad h_{\theta}(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x$$

$$\theta^T x = [\theta_0, \theta_1, \theta_2 \dots \theta_n] [x_0$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \dots \\ \theta_n \end{bmatrix} \quad x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \dots \\ x_n \end{bmatrix}$$

$$\begin{bmatrix} x_1 \\ x_2 \\ \dots \\ x_n \end{bmatrix} = \mathbb{R} = h_{\theta}(x)$$

a real number

Linear Regression

- How do we pick/calculate/learn the values for the θ_i 's?
 - As a starting point, we can make $h() \approx y$
- We can define a **cost function**:
 - $J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$
 - The superscript i 's refer to training examples, not raise to a power!
- We want to choose θ to minimize $J(\theta)$
- To do so, let's use a search algorithm that starts with some initial guess for θ , and repeatedly changes θ to make $J(\theta)$ smaller, until we converge at a minimum $J(\theta)$ value.

Linear Regression

- **Search algorithm = Gradient Descent**
 - Start with some initial guess at θ
 - and then repeatedly perform the update:
 - $\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta), j = 0 \dots n$
 - this rule is called the LMS update rule (least mean squares rule)
 - α (alpha) is the learning rate
 - Great care needs to be taken choosing alpha

Linear Regression

- **Gradient Descent**

- Calculating the derivative, $\frac{\partial}{\partial \theta_j} J(\theta)$ and for all n , we get:
- *repeat until convergence* {

$$\theta_j = \theta_j + \alpha \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)})) x_j^i, j = 0 \dots n$$

- }
- for each θ_j cycle through all m training examples
- This method is called **batch gradient descent**
- This $J(\theta)$ equation is a convex quadratic function, and has only a global minimum. Beware of functions that may have local minima, because gradient descent could get “stuck” in a local minima.

Linear Regression

- **Gradient Descent**

- Test for convergence is a **user defined function**
- Downside to **batch gradient descent** is that every θ update needs to run through all m training samples
 - Computationally expensive for large training sets

IIoT Intro



Source: World Economic Forum: <https://www.youtube.com/watch?v=8NGzrtK7eV0>

End